

DLDCALab 5: Stack Overflow

Week 6

September 4, 2026

Introduction

This lab moves from integers to floating point, and introduces real functions. The tasks build up from scalar floating-point arithmetic, through writing and composing your own functions under the **SystemV** calling convention, to recursion, along with an additional graded task.

Necessary tools

We shall be using `nasm` to assemble x86-64 assembly into object files, and `ld/gcc/g++` to link those object files into an executable. Before starting the lab, ensure both are installed:

```
$ nasm -v
NASM version [version]
$ ld -v # or gcc/g++ --version
GNU ld (GNU Binutils for Ubuntu) [version]
```

Structure of submission/templates

Submission format:

```
your-roll-no/
|- task1/
    |- dotprod.s (TODO)
|- task2/
    |- callee_caller.s (TODO)
    |- dump_regs.s
    |- to_a.s
|- task3/
    |- power.s (TODO)
|- labexam/
    |- poly_eval.s (TODO)
```

The folder structure of the expected submission for this lab is given above. We expect this folder to be compressed into a tarball with the naming convention of `your-roll-no.tar.gz` in lowercase. The command given below can be used to do the same

```
$ ./build_archive.sh
Enter Student Roll Number: your-roll-no
Packaging assignment files...
[+] Copied: ...
Successfully created your-roll-no.tar.gz
```

Background: Scalar Floating Point & Functions

x86-64 provides sixteen 128-bit `xmm` registers. For this lab, we only care about the low 64 bits of each register (enough to hold a `double`) and only the *scalar double* (`sd`) instruction forms, such as `movsd`, `addsd`, and `mulsd`. The *packed* forms (`ps/pd`) operate on several values at once and will be covered in the next lab.

As a reminder, the following facts hold true about the **SystemV** convention:

1. Integer/pointer arguments and floating-point arguments are assigned from two *separate* sequences of registers – the first six integer/pointer arguments go in `rdi`, `rsi`, `rdx`, `rcx`, `r8`, `r9`, while the first eight floating-point arguments go in `xmm0` through `xmm7`, independently of each other. A function taking a `float` and an `int` would put the `float` in `xmm0` and the `int` in `rdi`.
2. Return values follow the same split: an integer/pointer result comes back in `rax`, a floating-point result in `xmm0`.
3. Every `xmm` register is caller-saved. There is no floating-point equivalent of a callee-saved general-purpose register like `rbx/r1[2-5]`. If you need a value in an `xmm` register to survive a `call`, you are responsible for saving it yourself (the stack is the usual place).
4. Syscalls are like special functions. They go to the OS and tell it to perform some action. To utilize them, we put the ‘syscall number’ in `rax` and put arguments in `rdi`, `rsi`, `rdx`, `r10`, `r8`, `r9`.

Tasks

Task 1: Floating-Point Dot Product

The same dot product from your very first assembly lab, except `vec1` and `vec2` now hold 64-bit doubles instead of integers, and the arithmetic needs to happen in the `xmm` registers instead of the general-purpose ones. With the values provided, their dot product

$$\text{vec1}[0] \cdot \text{vec2}[0] + \text{vec1}[1] \cdot \text{vec2}[1] + \cdots + \text{vec1}[n-1] \cdot \text{vec2}[n-1]$$

comes out to 11.3. Since the exit code has to be an integer, truncate this to its integer part toward zero using `cvttsd2si` (round towards zero).

Fill in the body of `_start` in `dotprod.s`.

To run your code:

```
$ make task1
...
Return Code: 11
```

Task 2: Caller and Callee Semantics

Let's now write functions which are actually callable using `call` and return using `ret`. We will follow the `SystemV` convention throughout this lab. You are given `dump_regs()`, a function that prints every general-purpose register's current value and modifies nothing; calling it should look exactly like a no-op to the rest of your program. You'll use it to directly observe what a function call does and doesn't change.

(You can also use `gdb` instead of putting calls to `dump_regs` in your code.)

`func_callee`

Implement `func_callee(int a{rdi}, int b{rsi}) -> rax`, which returns `a + b`. Once you've computed the sum, deliberately overwrite every *other* caller-saved register (`rcx`, `rdx`, `rsi`, `rdi`, `r8`, `r9`, `r10`, `r11`) with any junk value you like. This models a "worst case" callee – one that gives you no guarantees beyond what the ABI actually promises.

`func_caller`

Implement `func_caller() -> rax`. It needs to call `func_callee(10, 20)`, but it also needs a multiplier of 3 to survive that call, so that it can return `func_callee`'s result times 3 afterward.

Call `dump_regs()` once right before calling `func_callee`, and once right after it returns, so you have a before-and-after snapshot of every register to compare.

Putting it together

In `_start`, call `func_caller()` and exit with its result.

Before you move on: look at the two `dump_regs` outputs side by side. Which registers changed, and which didn't? For each one, why? Is it because the ABI guarantees it? Is it because `func_callee` happened to leave it alone anyway? Or is it something else? In particular, check what happened to whichever register you chose to keep your multiplier in.

Fill in `func_callee`, `func_caller`, and `_start` in `callee_caller.s`.

To run your code:

```
$ make task2
rax: 0    rbx: 3    rcx: 0    ...
rax: 30   rbx: 3    rcx: 57005 ...
Return Code: 90
```

Task 3: Recursive Power, Two Ways

You're given a double `base` and a non-negative integer `exponent` in `.data`. This task asks for two different implementations of the same function, so you can see the difference between ordinary recursion and a tail call directly.

3a: Naive recursion

Implement `pow_naive(double base{xmm0}, uint64_t exponent{rdi}) -> xmm0` using *exponentiation by squaring*, i.e. the recurrence

$$b^e = \begin{cases} 1 & e = 0 \\ (b \times b)^{e/2} & e \text{ even} \\ b \times (b \times b)^{(e-1)/2} & e \text{ odd} \end{cases}$$

using ordinary `call/ret` recursion.

A reminder about SystemV's conventions:

1. `base` is a double, so it's passed in `xmm0`, the first *floating-point* argument register. `exponent` is an integer, so it's passed in `rdi`, the first *integer* argument register. Register locations for `double pow(double, uint64_t)` and `double pow(uint64_t, double)` will be identical.
2. `xmm0` does not survive a `call` unless you save it first. The odd case above needs the original `base` *after* the recursive call returns, by which point `xmm0` has been reused for the call's argument and return value.

3b: Tail-call rewrite

Implement `pow_tail(double base{xmm0}, uint64_t exponent{rdi}, double acc{xmm1}=1.0)` -> `xmm0`, computing the exact same result, but rewritten such that we use `jmp` to go back to the function. Thread an accumulator through `xmm1` (call it initially with `acc = 1.0`) so that the multiplication which used to happen *after* the recursive call in 3a now happens *before* it instead.

Implement this using `jmp` instead of `call/ret` for the recursive step. To see what changes, step through a run of your program and see the value of `rsp`. What do you observe?

Putting it together

In `_start`, call both `pow_naive(base, exponent)` and `pow_tail(base, exponent, 1.0)`, truncate each result toward zero, and exit with their sum. Since both should independently come out to 57, a correct solution exits with 114.

Fill in `pow_naive`, `pow_tail`, and `_start` in `power.s`.

To run your code:

```
$ make task3
...
Return Code: 114
```

Lab Exam: Polynomial Evaluation

You are given an array of doubles, `coeffs`, its length `coeffs_len`, and a double `x_val`. Treating `coeffs` as the coefficients of a polynomial (`coeffs[0]` is the constant term, `coeffs[1]` the coefficient of x^1 , and so on), evaluate

$$\text{coeffs}[0] \cdot x^0 + \text{coeffs}[1] \cdot x^1 + \dots + \text{coeffs}[n-1] \cdot x^{n-1}$$

at $x = x_val$, truncate the result toward zero, and exit with that as the exit code.

You'll implement this recursively, using Horner's method – which just means factoring x out repeatedly:

$$c_0 + x(c_1 + x(c_2 + \dots))$$

so that each step only needs one multiply and one add. As a recursive function over "the coefficients from some starting point onward", this is:

```
poly_eval(coeffs, n, x):  
    if n == 1:  
        return coeffs[0]  
  
    rest = poly_eval(coeffs + 8, n - 1, x)  
    return coeffs[0] + x * rest
```

Implement this as `poly_eval(double* coeffs{rdi}, size_t N{rsi}, x{xmm0})`.

Fill in `poly_eval` and `_start` in `poly_eval.s`.

To run your code:

```
$ make labexam  
...  
Return Code: 49
```

Table 1: Some basic x86-64 Assembly Instructions

Instruction	Operands	Effect
mov	reg, reg/*mem/imm	$\text{reg} \leftarrow \text{reg}/\text{*mem}/\text{imm}$
mov	*mem, reg	$\text{*mem} \leftarrow \text{reg}$
movzx / movsx	reg, reg/*mem	$\text{reg} \leftarrow \text{reg}/\text{*mem}, \text{zero-/sign-extended}$
lea	reg, *mem	$\text{reg} \leftarrow \text{address of *mem}$
add	reg, reg/*mem/imm	$\text{reg} \leftarrow \text{reg} + \text{reg}/\text{*mem}/\text{imm}$
add	*mem, reg/imm	$\text{*mem} \leftarrow \text{*mem} + \text{reg}/\text{imm}$
sub	reg, reg/*mem/imm	$\text{reg} \leftarrow \text{reg} - \text{reg}/\text{*mem}/\text{imm}$
sub	*mem, reg/imm	$\text{*mem} \leftarrow \text{*mem} - \text{reg}/\text{imm}$
inc	reg/*mem	$\text{reg}/\text{*mem} \leftarrow \text{reg}/\text{*mem} + 1$
dec	reg/*mem	$\text{reg}/\text{*mem} \leftarrow \text{reg}/\text{*mem} - 1$
neg	reg/*mem	$\text{reg}/\text{*mem} \leftarrow -\text{reg}/\text{*mem}$ (Two's complement)
not	reg/*mem	$\text{reg}/\text{*mem} \leftarrow \neg \text{reg}/\text{*mem}$ (One's complement/bitwise)
shl / shr	reg/*mem, imm/cl	Logical shift left / right by imm or cl bits
sar	reg/*mem, imm/cl	Arithmetic shift right (preserves sign bit)
imul	reg, reg/*mem/imm	$\text{reg} \leftarrow \text{reg} \times \text{reg}/\text{*mem}/\text{imm}$
imul	*mem, reg/imm	$\text{*mem} \leftarrow \text{*mem} \times \text{reg}/\text{imm}$
div	reg	$\text{rax} \leftarrow (\text{rdx}:\text{rax})/\text{reg}$ $\text{rdx} \leftarrow (\text{rdx}:\text{rax})\% \text{reg}$
and	reg, reg/*mem/imm	$\text{reg} \leftarrow \text{reg} \wedge \text{reg}/\text{*mem}/\text{imm}$
and	*mem, reg/imm	$\text{*mem} \leftarrow \text{*mem} \wedge \text{reg}/\text{imm}$
or	reg, reg/*mem/imm	$\text{reg} \leftarrow \text{reg} \vee \text{reg}/\text{*mem}/\text{imm}$
or	*mem, reg/imm	$\text{*mem} \leftarrow \text{*mem} \vee \text{reg}/\text{imm}$
xor	reg, reg/*mem/imm	$\text{reg} \leftarrow \text{reg} \oplus \text{reg}/\text{*mem}/\text{imm}$
xor	*mem, reg/imm	$\text{*mem} \leftarrow \text{*mem} \oplus \text{reg}/\text{imm}$
cmp	reg/*mem, reg/imm	Sets flags based on subtraction: $\text{reg}/\text{*mem} - \text{reg}/\text{imm}$. ZF = 1 if equal, SF = 1 if result negative, CF = 1 if borrow (unsigned), OF = 1 if signed overflow.
test	reg/*mem, reg/imm	Sets flags based on bitwise AND: $\text{reg}/\text{*mem} \wedge \text{reg}/\text{imm}$. ZF = 1 if result is zero, SF = 1 if high bit set, CF and OF are cleared. No result is stored.
jmp	label	Unconditional jump to label
je / jz	label	Jump if Zero Flag (ZF) = 1
jne / jnz	label	Jump if Zero Flag (ZF) = 0
jl	label	Jump if Less (SF $\neq$ OF)
jle	label	Jump if Less or Equal (ZF = 1 or SF $\neq$ OF)
jg	label	Jump if Greater (ZF = 0 and SF = OF)
jge	label	Jump if Greater or Equal (SF = OF)
cmovcc	reg, reg/*mem	$\text{reg} \leftarrow \text{reg}/\text{*mem}$ if cc holds, else unchanged.
syscall		Performs a syscall with syscall number stored in rax Arguments passed in other registers starting from rdi
push	reg	Pushes the value in reg onto the stack
pop	reg	Pops the value on top of the stack onto reg
call	label	Calls the function called label
ret		Pops the return address from stack and jumps to it

Table 2: Scalar Floating-Point (`xmm`) Instructions

Instruction	Operands	Effect
<code>movsd</code>	<code>xmm, xmm/*mem</code>	$xmm \leftarrow xmm/*mem$ (low 64 bits)
<code>movsd</code>	<code>*mem, xmm</code>	$*mem \leftarrow xmm$ (low 64 bits)
<code>movq</code>	<code>xmm, reg/*mem</code>	$xmm \leftarrow reg/*mem$ (raw 64 bits, no conversion)
<code>movq</code>	<code>reg/*mem, xmm</code>	$reg/*mem \leftarrow xmm$ (raw 64 bits, no conversion)
<code>xorps / xorpd</code>	<code>xmm, xmm</code>	$xmm \leftarrow xmm \oplus xmm$ (idiom to zero an <code>xmm</code> register)
<code>pxor</code>	<code>xmm, xmm</code>	$xmm \leftarrow xmm \oplus xmm$ (bitwise XOR on <code>xmm</code>)
<code>addsd</code>	<code>xmm, xmm/*mem</code>	$xmm \leftarrow xmm + xmm/*mem$
<code>subsd</code>	<code>xmm, xmm/*mem</code>	$xmm \leftarrow xmm - xmm/*mem$
<code>mulsd</code>	<code>xmm, xmm/*mem</code>	$xmm \leftarrow xmm \times xmm/*mem$
<code>divsd</code>	<code>xmm, xmm/*mem</code>	$xmm \leftarrow xmm \div xmm/*mem$
<code>comisd</code>	<code>xmm, xmm/*mem</code>	Sets flags like <code>cmp</code> , comparing as doubles (use <code>ja/jae/jb/jbe</code>). Raises exception on NaN.
<code>ucomisd</code>	<code>xmm, xmm/*mem</code>	Same as <code>comisd</code> , but does not raise exception on NaN
<code>cvttsd2si</code>	<code>reg, xmm/*mem</code>	$reg \leftarrow$ integer part, truncated toward zero
<code>cvtsi2sd</code> ¹	<code>xmm, reg/*mem</code>	$xmm \leftarrow reg/*mem$ converted to a double

¹Prefer loading floating point literals or the binary representation directly instead of using this. Use this only for runtime `int` to `double` conversions